

In [1]: ##### PART 1 - Installation and Intial Setup #####

VBox()

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
2	application_1682087553656_0003	pyspark	idle	Link	Link	✓

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

SparkSession available as 'spark'.

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

In [2]: %%info

Current session configs: {'conf': {'spark.pyspark.python': 'python3',

'spark.pyspark.virtualenv.enabled': 'true',

'spark.pyspark.virtualenv.type': 'native',

'spark.pyspark.virtualenv.bin.path': '/usr/bin/virtualenv'}, 'kind':

'pyspark'}

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
2	application_1682087553656_0003	pyspark	idle	Link	Link	✓

In [3]: sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")

VBox()

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

Collecting pandas==1.0.3

Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85edledc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)

Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)

Using cached https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bbb257d0355c7f6128853c78955f57342a56d/python_dateutil-2.8.2-py2.py3-none-any.whl

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)

Installing collected packages: python-dateutil, pandas

Successfully installed pandas-1.0.3 python-dateutil-2.8.2

Collecting matplotlib==3.2.1

Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/6c/10/a7d0fa5baea8fe7b50f448ab742f26f52b80bfca85ac2be9d35cdd9a3246/pyparsing-3.0.9-py3-none-any.whl>

Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1682096917141-0/lib/python3.7/site-packages (from matplotlib==3.2.1)

Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)

Collecting cyclor>=0.10 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f51/cyclor-0.11.0-py3-none-any.whl>

Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/ab/8f/8dbe2d4efc4c0b08ec67d6efb7cc31fbfd688c80afad85f65980633b0d37/kiwisolver-1.4.4-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)

Collecting typing-extensions; python_version < "3.8" (from kiwisolver>=1.0.1->matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/31/25/5abcd82372d3d4a3932elfa8c3dbf9efac10cc7c0d16e78467460571b404/typing_extensions-4.5.0-py3-none-any.whl

Installing collected packages: pyparsing, cyclor, typing-extensions, kiwisolver, matplotlib

Successfully installed cyclor-0.11.0 kiwisolver-1.4.4 matplotlib-3.2.1 pyparsing-3.0.9 typing-extensions-4.5.0

```
In [4]: # The data comes from https://www.kaggle.com/datasets/ashirwadsangwan/imdb-data
actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/name.basics.tsv',
genres = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.basics.tsv',
movie_actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.princi
movie_ratings = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.ratir
```

VBox()

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
In [5]: ## ACTORS
```

VBox()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [6]: actors.printSchema()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
root
|-- nconst: string (nullable = true)
|-- primaryName: string (nullable = true)
|-- birthYear: string (nullable = true)
|-- deathYear: string (nullable = true)
|-- primaryProfession: string (nullable = true)
|-- knownForTitles: string (nullable = true)
```

```
In [7]: actors.select("primaryName", "birthYear", "deathYear", "knownForTitles").show(5)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+-----+-----+
| primaryName | birthYear | deathYear | knownForTitles |
+-----+-----+-----+-----+
| Fred Astaire | 1899 | 1987 | tt0050419,tt00531... |
| Lauren Bacall | 1924 | 2014 | tt0071877,tt01170... |
| Brigitte Bardot | 1934 | \N | tt0054452,tt00491... |
| John Belushi | 1949 | 1982 | tt0077975,tt00725... |
| Ingmar Bergman | 1918 | 2007 | tt0069467,tt00509... |
+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [8]: ## GENRES
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
only showing top 5 rows
```

```
In [9]: genres.select("titleType", "primaryTitle", "genres").show(10)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+-----+
| titleType | primaryTitle | genres |
+-----+-----+-----+
| short | Carmencita | Documentary,Short |
| short | Le clown et ses c... | Animation,Short |
| short | Pauvre Pierrot | Animation,Comedy,... |
| short | Un bon bock | Animation,Short |
| short | Blacksmith Scene | Comedy,Short |
| short | Chinese Opium Den | Short |
| short | Corbett and Court... | Short,Sport |
| short | Edison Kinetoscop... | Documentary,Short |
| movie | Miss Jerry | Romance |
| short | Exiting the Factory | Documentary,Short |
+-----+-----+-----+
only showing top 10 rows
```

```
In [10]: genres.select("titleType").distinct().show()
```

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+
|   titleType   |
+-----+
|   tvSeries    |
| tvMiniSeries  |
|   movie       |
|   videoGame   |
|   tvSpecial   |
|   video       |
|   tvMovie     |
|   tvEpisode   |
|   tvShort     |
|   short       |
+-----+

```

In [11]: `genres.printSchema()`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
root
|-- tconst: string (nullable = true)
|-- titleType: string (nullable = true)
|-- primaryTitle: string (nullable = true)
|-- originalTitle: string (nullable = true)
|-- isAdult: string (nullable = true)
|-- startYear: string (nullable = true)
|-- endYear: string (nullable = true)
|-- runtimeMinutes: string (nullable = true)
|-- genres: string (nullable = true)

```

In [12]: `#MOVIE ACTORS`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

In [13]: `movie_actors.printSchema()`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
root
|-- tconst: string (nullable = true)
|-- ordering: string (nullable = true)
|-- nconst: string (nullable = true)
|-- category: string (nullable = true)
|-- job: string (nullable = true)
|-- characters: string (nullable = true)

```

In [14]: `movie_actors.show(10)`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

```

+-----+-----+-----+-----+-----+-----+
+
|  tconst|ordering|  nconst|      category|      job| characters
|
+-----+-----+-----+-----+-----+-----+
+
|tt0000001|      1|nm1588970|      self|      \N|["Herself"]
|tt0000001|      2|nm0005690|    director|      \N|      \N
|tt0000001|      3|nm0374658|cinematographer|director of photo...|      \N
|tt0000002|      1|nm0721526|    director|      \N|      \N
|tt0000002|      2|nm1335271|    composer|      \N|      \N
|tt0000003|      1|nm0721526|    director|      \N|      \N
|tt0000003|      2|nm5442194|    producer|    producer|      \N
|tt0000003|      3|nm1335271|    composer|      \N|      \N
|tt0000003|      4|nm5442200|    editor|      \N|      \N
|tt0000004|      1|nm0721526|    director|      \N|      \N
|
+-----+-----+-----+-----+-----+-----+
+
only showing top 10 rows

```

In [15]: `#MOVIE RATINGS`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

In [16]: `movie_ratings.printSchema()`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
root
|-- tconst: string (nullable = true)
|-- averageRating: string (nullable = true)
|-- numVotes: string (nullable = true)

```

In [17]: `movie_ratings.show(10)`

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

tconst	averageRating	numVotes
tt0000001	5.6	1550
tt0000002	6.1	186
tt0000003	6.5	1207
tt0000004	6.2	113
tt0000005	6.1	1934
tt0000006	5.2	102
tt0000007	5.5	615
tt0000008	5.4	1667
tt0000009	5.4	81
tt0000010	6.9	5545

only showing top 10 rows

In [18]: *#OVERVIEW of Data*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [19]: *# Display number of rows and columns in the actors DataFrame*

```
print("Number of rows: ", actors.count())
print("Number of columns: ", len(actors.columns))

# Display number of rows and columns in the genres DataFrame
print("Number of rows: ", genres.count())
print("Number of columns: ", len(genres.columns))

# Display number of rows and columns in the movie_actors DataFrame
print("Number of rows: ", movie_actors.count())
print("Number of columns: ", len(movie_actors.columns))

# Display number of rows and columns in the movie_ratings DataFrame
print("Number of rows: ", movie_ratings.count())
print("Number of columns: ", len(movie_ratings.columns))
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
Number of rows: 9706922
Number of columns: 6
Number of rows: 6321302
Number of columns: 9
Number of rows: 36468817
Number of columns: 6
Number of rows: 993153
Number of columns: 3
```

In [20]: *##### PART 2 - Analyzing Genres #####*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [21]: `genres.select("tconst","titleType","genres").show(5)`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
+-----+-----+-----+
| tconst|titleType| genres|
+-----+-----+-----+
|tt0000001| short| Documentary,Short|
|tt0000002| short| Animation,Short|
|tt0000003| short| Animation,Comedy,...|
|tt0000004| short| Animation,Short|
|tt0000005| short| Comedy,Short|
+-----+-----+-----+
only showing top 5 rows
```

```
In [22]: from pyspark.sql.functions import split, explode

# Split genres by comma and create a new column 'genre'
genre_split = genres.withColumn('genre', explode(split(genres.genres, ',')))

# Display the resulting table
genre_split.select('tconst', 'titleType', 'genre').show(10)

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
+-----+-----+-----+
| tconst|titleType| genre|
+-----+-----+-----+
|tt0000001| short| Documentary|
|tt0000001| short| Short|
|tt0000002| short| Animation|
|tt0000002| short| Short|
|tt0000003| short| Animation|
|tt0000003| short| Comedy|
|tt0000003| short| Romance|
|tt0000004| short| Animation|
|tt0000004| short| Short|
|tt0000005| short| Comedy|
+-----+-----+-----+
only showing top 10 rows
```

```
In [23]: ##### Total Unique Genres #####

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
In [24]: genre_split = genres.withColumn('genre', explode(split(genres.genres, ',')))

# Get the distinct genres and count them
total_genres = genre_split.select("genre").distinct().count()

# Print the result
print("Total number of unique genres in the movie category: ", total_genres)

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
Total number of unique genres in the movie category: 29
```

```
In [25]: #What are the unique genres available?

VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
In [26]: unique_genres = genre_split.select("genre").distinct()
unique_genres.show()
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
+-----+
|      genre      |
+-----+
|      Mystery    |
|      Musical    |
|      Sport      |
|      Action     |
|      Talk-Show  |
|      Romance    |
|      Thriller   |
|      \N         |
|      Reality-TV |
|      Family     |
|      Fantasy    |
|      History    |
|      Animation  |
|      Short      |
|      Film-Noir  |
|      Sci-Fi     |
|      News       |
|      Drama      |
|      Documentary|
|      Western    |
+-----+
```

only showing top 20 rows

```
In [27]: from pyspark.sql.functions import col

# Get the distinct genres and filter out null values
unique_genres = genre_split.select("genre").distinct().filter(col("genre") != '

# Show the unique genres
unique_genres.show()
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```



```
+-----+
|      genre      |
+-----+
|      Mystery    |
|      Musical    |
|      Sport      |
|      Action     |
|      Talk-Show  |
|      Romance    |
|      Thriller   |
|      Reality-TV |
|      Family     |
|      Fantasy    |
|      History    |
|      Animation  |
|      Film-Noir  |
|      Short      |
|      Sci-Fi     |
|      News       |
|      Drama      |
|      Documentary|
|      Western    |
|      Comedy     |
+-----+
```

only showing top 20 rows

```
In [28]: total_genres = unique_genres.select("genre").distinct().count()

# Print the result
print("Total number of unique genres in the movie category: ", total_genres)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
Total number of unique genres in the movie category: 28
```

```
In [29]: ##### TOP GENRES BY MOVIES #####
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

```
In [30]: from pyspark.sql.functions import col
genre_split_filtered = genre_split.filter(col('titleType') == 'movie')

filtered_df = genre_split_filtered.filter(col('genre') != '\\N')

joined_dataframe = movie_ratings.join(filtered_df, "tconst", "inner")

joined_dataframe.select('genre', 'averageRating').show(10)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

genre	averageRating
Drama	4.2
Drama	4.2
Biography	4.1
Drama	4.1
History	4.1
Drama	5.7
Drama	4.6
History	4.6
Biography	6.3
Drama	6.3

only showing top 10 rows

```
In [31]: from pyspark.sql.functions import avg

avg_rating_by_genre = (
    joined_dataframe
    .select(col("genre"), col("averageRating").cast("float").alias("averageRating"))
    .groupby("genre")
    .agg(avg("averageRating").alias("avg_rating"))
)
avg_rating_by_genre.show(20)
```

VBox()

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

genre	avg_rating
Mystery	5.940437537126316
Musical	6.203246053185319
Action	5.718734067904495
Sport	6.600145190943391
Talk-Show	5.800000190734863
Romance	6.125714179294426
Thriller	5.625967567519544
Reality-TV	6.379310377712907
Family	6.250560452699635
Fantasy	5.924820762891499
History	6.822718117193864
Animation	6.326203749467441
Film-Noir	6.636246780503378
Short	7.259999942779541
Sci-Fi	5.325150006900168
News	7.200916040944689
Drama	6.288080211097538
Documentary	7.245469805371099
Western	5.948970991005059
Comedy	5.941363107822231

only showing top 20 rows

```
In [32]: ### Horizontal Bar Chart of Top Genres

#With this data available, let us now build a barchart of all genres

##*HINT*: don't forget about the matplotlib magic!
```

```
##matplotlib plt
```

```
VBox()  
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [33]: **from** pyspark.sql.functions **import** desc

```
genre_avg_rating = avg_rating_by_genre.orderBy(desc('avg_rating'))  
genre_avg_rating.show(20)
```

```
VBox()  
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
+-----+-----+  
|      genre|      avg_rating|  
+-----+-----+  
|      Short|7.259999942779541|  
|Documentary|7.245469805371099|  
|      News|7.200916040944689|  
|Biography|6.983637643044585|  
|Game-Show|6.974999904632568|  
|    History|6.822718117193864|  
|    Music|6.752020207214588|  
|Film-Noir|6.636246780503378|  
|    Sport|6.600145190943391|  
|    War|6.483807036278403|  
|Reality-TV|6.379310377712907|  
|Animation|6.326203749467441|  
|    Drama|6.288080211097538|  
|    Family|6.250560452699635|  
|Musical|6.203246053185319|  
|Romance|6.125714179294426|  
|    Crime|6.026013333109149|  
|Western|5.948970991005059|  
|    Comedy|5.941363107822231|  
|Mystery|5.940437537126316|  
+-----+-----+
```

only showing top 20 rows

In [34]: **import** pandas **as** pd
import matplotlib.pyplot **as** plt
from pyspark.sql.functions **import** asc

```
genre_avg_rating = avg_rating_by_genre.orderBy(asc('avg_rating'))  
genre_ratings_pd = genre_avg_rating.toPandas()
```

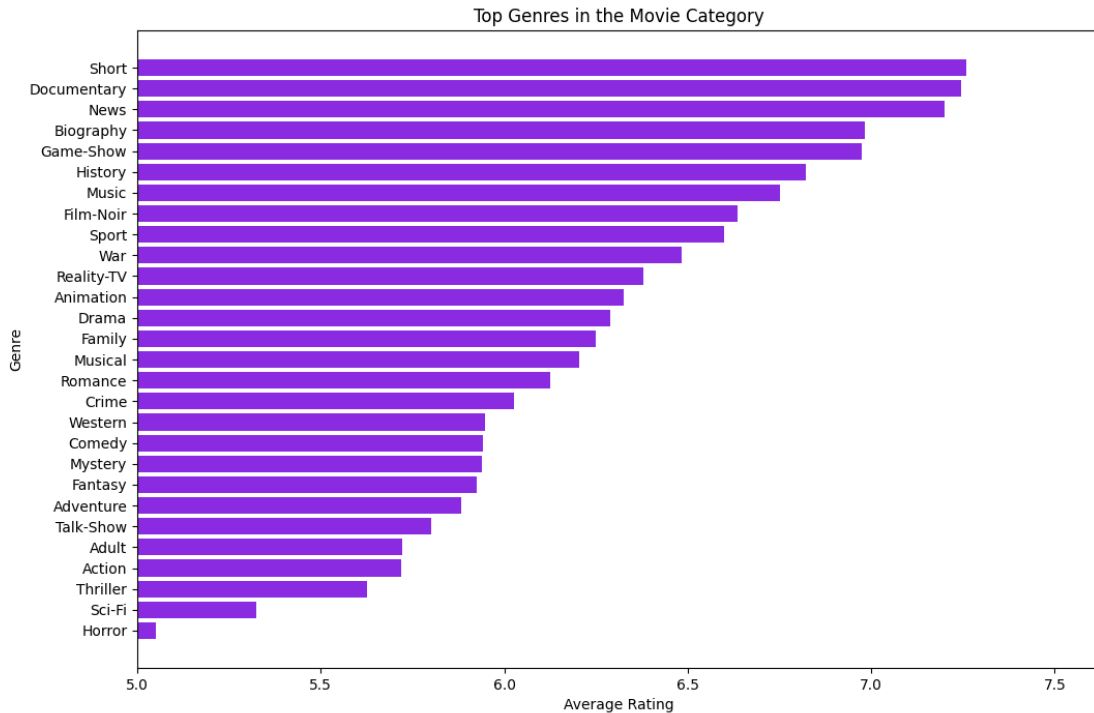
```
# Plot the data using Matplotlib  
fig, ax = plt.subplots(figsize=(12, 8))
```

```
ax.barh(genre_ratings_pd["genre"], genre_ratings_pd["avg_rating"], color="bluev  
ax.set_title("Top Genres in the Movie Category")  
ax.set_xlabel("Average Rating")  
ax.set_ylabel("Genre")
```

```
# Set the minimum value of the x-axis to 5  
ax.set_xlim([5.0, None])
```

```
# Display the plot using the %matplotlib magic command  
%matplotlib plt
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```



In [35]: *## PART 3 - Analyzing Job Categories*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [36]: *## Total Unique Job Categories*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [37]: *### What is the total number of unique job categories? ###*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [38]: `movie_actors.printSchema()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
root
|-- tconst: string (nullable = true)
|-- ordering: string (nullable = true)
|-- nconst: string (nullable = true)
|-- category: string (nullable = true)
|-- job: string (nullable = true)
|-- characters: string (nullable = true)
```

In [39]: `movie_actors.select('tconst','category').show(5)`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+
| tconst | category |
+-----+-----+
| tt0000001 | self |
| tt0000001 | director |
| tt0000001 | cinematographer |
| tt0000002 | director |
| tt0000002 | composer |
+-----+-----+
only showing top 5 rows
```

In [61]: `from pyspark.sql.functions import countDistinct`

```
num_categories = movie_actors.select(countDistinct('category')).collect()[0][0]
print('The total number of unique job categories is:', num_categories)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
The total number of unique job categories is: 12
```

In [41]: `unique_category = movie_actors.select('category').distinct()
unique_category.show()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+
| category |
+-----+-----+
| actress |
| producer |
| production_designer |
| writer |
| actor |
| cinematographer |
| archive_sound |
| archive_footage |
| self |
| editor |
| composer |
| director |
+-----+-----+
```

In [42]: `## Top Job Categories`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

In [43]: `### Counts of Titles / Job Category`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)

```

```
In [44]: from pyspark.sql.functions import count

category_count = movie_actors.groupBy('category').agg(count('*').alias('count'))

category_count.show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+
|          category|    count|
+-----+-----+
|          actress| 6325097|
|          producer| 2197866|
|production_designer| 285924|
|           writer| 4811596|
|           actor| 8493701|
|cinematographer| 1300404|
|   archive_sound|    2143|
|archive_footage| 209035|
|           self| 6153089|
|           editor| 1197669|
|          composer| 1313187|
|          director| 4179106|
+-----+-----+
```

```
In [45]: ### Bar Chart of Top Job Categories

#With this data available, let us now build a barchart of the top 5 categories.

#HINT: don't forget about the matplotlib magic!

##matplotlib plt
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+
|          category|    count|
+-----+-----+
|          actress| 6325097|
|          producer| 2197866|
|production_designer| 285924|
|           writer| 4811596|
|           actor| 8493701|
|cinematographer| 1300404|
|   archive_sound|    2143|
|archive_footage| 209035|
|           self| 6153089|
|           editor| 1197669|
|          composer| 1313187|
|          director| 4179106|
+-----+-----+
```

```
In [46]: category_count.orderBy(desc('count')).show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+-----+
|          category|    count|
+-----+-----+
|          actress| 6325097|
|          producer| 2197866|
|production_designer| 285924|
|           writer| 4811596|
|           actor| 8493701|
|cinematographer| 1300404|
|   archive_sound|    2143|
|archive_footage| 209035|
|           self| 6153089|
|           editor| 1197669|
|          composer| 1313187|
|          director| 4179106|
+-----+-----+
```

```
In [47]: import pandas as pd
import matplotlib.pyplot as plt
from pyspark.sql.functions import desc

category_count_pd = category_count.orderBy(desc('count')).limit(5).toPandas()

# Set the width of the bars
bar_width = 0.5

# Plot the data using Matplotlib
fig, ax = plt.subplots(figsize=(15,9))

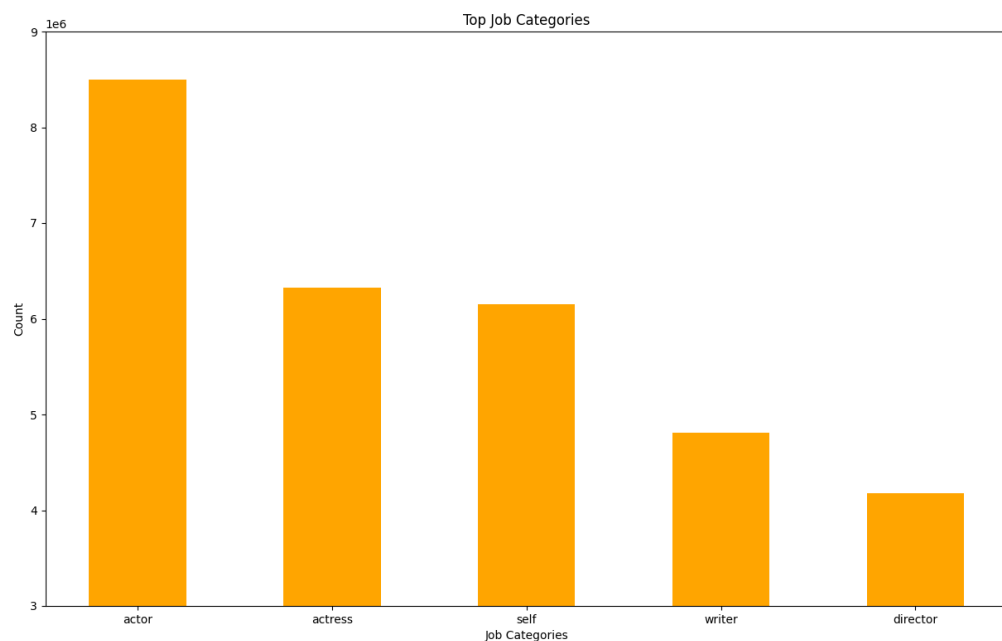
ax.bar(category_count_pd["category"], category_count_pd["count"], color="orange")
ax.set_title("Top Job Categories")
ax.set_xlabel("Job Categories")
ax.set_ylabel("Count")

# Set the minimum value of the y-axis to 3m and max to 9m
ax.set_ylim([3000000, 9000000])

# Display the plot using the %matplotlib magic command
%matplotlib plt
```

VBox()

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...



```
In [48]: # PART 4 - Answer to the following questions:
```

VBox()

FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
In [49]: ## 1) Find all the "movies" featuring "Johnny Depp" and "Helena Bonham Carter".
```

VBox()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [50]: ##### 1)

#joining tables to do query
actors_genres = actors.join(movie_actors,actors.nconst == movie_actors.nconst,
joined_actors_genres = actors_genres.join(genres,genres.tconst==actors_genres.tconst)

#Finding result
movies = joined_actors_genres.filter((joined_actors_genres.primaryName=="Johnny Depp"))
movies = movies.filter(movies.titleType=="movie").groupBy('tconst','primaryTitle').select('primaryTitle').show(truncate=False)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...)
+-----+
|primaryTitle|
+-----+
|Dark Shadows|
|Sweeney Todd: The Demon Barber of Fleet Street|
|Alice Through the Looking Glass|
|Alice in Wonderland|
|Charlie and the Chocolate Factory|
|Corpse Bride|
+-----+
```

```
In [51]: ## 2) Find all the "movies" featuring "Brad Pitt" after 2010.
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [52]: #2)
# Finding result
bradpitt = joined_actors_genres.filter((joined_actors_genres.primaryName=="Brad Pitt"))
bradpitt = bradpitt.filter(bradpitt.titleType == 'movie').filter(bradpitt.startYear > 2010)
BP_after2010_movies = bradpitt.orderBy(desc('startYear'))
BP_after2010_movies.show(truncate=False)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```


primaryTitle	startYear
Babylon	2021
Irresistible	2020
Kajillionaire	2020
Once Upon a Time ... in Hollywood	2019
Ad Astra	2019
The King	2019
Vice	2018
War Machine	2017
Allied	2016
Voyage of Time: Life's Journey	2016
The Big Short	2015
Hitting the Apex	2015
By the Sea	2015
Fury	2014
12 Years a Slave	2013
Kick-Ass 2	2013
World War Z	2013
Killing Them Softly	2012
The Tree of Life	2011
Moneyball	2011

In [53]: *## 3) What is the number of "movies" "acted" by "Zendaya" per year?*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [54]: *#3)*
Joining tables to do query
 actors_genres = actors.join(movie_actors,actors.nconst == movie_actors.nconst,
 joined_actors_genres = actors_genres.join(genres,genres.tconst==actors_genres.tconst)

```
#Finding result
zendaya = joined_actors_genres.filter((joined_actors_genres.primaryName=="Zendaya"))
zendaya = zendaya.filter(zendaya.category=="actress").select('startYear')
zendaya = zendaya.filter(zendaya.startYear.endswith('N')==False)
zendaya.groupby('startYear').count().show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

startYear	count
2020	1
2018	2
2017	1

In [55]: *##### 4) What are the "movies" by average rating greater than "9.7" and released after 2010?*

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

In [56]: **from** pyspark.sql.functions **import** col, when

```
# Load the title.basics dataset and filter only the movies released in 2019
movies_2019 = genres.filter((col("titleType") == "movie") & (col("startYear") =

# Load the title.ratings dataset and filter only the movies with an average rat
top_rated_movies = movie_ratings.filter(col("averageRating") > 9.7)

# Join the two datasets on the tconst column
joined_data = movies_2019.join(top_rated_movies, "tconst")

# Create a new column that maps averageRating to a numeric value for sorting
joined_data = joined_data.withColumn("rating_value", when(col("averageRating")
joined_data = joined_data.withColumn("rating_value", when(col("averageRating")
joined_data = joined_data.withColumn("rating_value", when(col("averageRating")

# Show the resulting dataframe with the required columns, sorted by rating_val
joined_data = joined_data.select("primaryTitle", "averageRating").orderBy(col('
joined_data.show(truncate=False)
```

VBox ()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

primaryTitle	averageRating
A Grunt's Life	10.0
Kirket	10.0
Bu Can Var Oldugu Sürece	10.0
L'Enfant Terrible	10.0
The Butcher Baronet	10.0
A Medicine for the Mind	10.0
Our Scripted Life	10.0
Love in Kilnerry	10.0
The Twilight Zone: A 60th Anniversary Celebration	10.0
Superhombre	9.9
The Cardinal	9.9
Puritan: All of Life to The Glory of God	9.9
Kamen Rider Zi-O: Over Quartzer	9.8
Time and motion	9.8
We Shall Not Die Now	9.8
Gini Helida Kathe	9.8
Square One	9.8
From Shock to Awe	9.8
Randhawa	9.8

```
In [57]: ## Extra Credit - Analysis of your choice
```

```
## Try and analyze some interesting dimension to this data. You should specify
### You must join at least two datasets.
```

VBox ()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
In [58]: ## For this section. We'll answer this question: "Highest Rated Movies Featurin
```

VBox ()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
In [59]: actors_genres = actors.join(movie_actors,actors.nconst == movie_actors.nconst,
joined_actors_genres = actors_genres.join(genres,genres.tconst==actors_genres.tconst)
extra_credit_df = joined_actors_genres.join(movie_ratings,movie_ratings.tconst

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

```
In [60]: from pyspark.sql.functions import desc

leodicaprio_movies = extra_credit_df.filter(extra_credit_df.primaryName == "Leo
        .filter(extra_credit_df.titleType == "movie") \
        .filter(extra_credit_df.averageRating >= 8.0) \
        .select("primaryTitle", "averageRating") \
        .orderBy(desc("averageRating"))

leodicaprio_movies.show(truncate=False)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Lay
out(height='25px', width='50%'),...
```

primaryTitle	averageRating
Inception	8.8
The Clock	8.8
The Departed	8.5
Django Unchained	8.4
Before the Flood	8.3
The Wolf of Wall Street	8.2
Shutter Island	8.1
Catch Me If You Can	8.1
Struggle: The Life and Lost Art of Szukalski	8.0
The Revenant	8.0
Once Upon a Time ... in Hollywood	8.0
Blood Diamond	8.0